

A Comprehensive Benchmark of Deep Learning Architectures for 2D Echocardiographic Segmentation: From CNNs to Transformers

Fares Sofiane Mazouni^{a,*}, Mohamed Yacine Haouam^a, Mohamed Amroune^a, Issam Bendib^a

^a*Laboratory of Mathematics, Informatics and Systems (LAMIS), Echahid Cheikh Larbi Tebessi University, Tebessa,
12002, Algeria*

Abstract

Two-dimensional echocardiography is the most widely used cardiac imaging modality, but quantitative analysis is hampered by speckle noise, low boundary contrast, and inter-observer variability. U-Net descendants, attention-augmented CNNs, hybrid CNN–Transformer designs, and pure Transformer networks all report strong results on this task, yet direct comparison between them is confounded by inconsistent training protocols and evaluation pipelines. We benchmark nine representative architectures spanning three paradigms under one training protocol, one data split, and one evaluation pipeline on the CAMUS dataset, reporting per-class Dice and IoU, boundary distances in millimetres, biplane Simpson’s ejection fraction (EF) error, phase-stratified and image-quality-stratified results, efficiency, and Bonferroni-corrected Wilcoxon tests. Two findings stand out. First, boundary accuracy has saturated: the best architecture reaches $HD_{95} = 4.08$ mm against a published CAMUS baseline of ≈ 4.3 mm, and the top four architectures lie within 0.2 mm of one another, well inside training variance. Clinical agreement, by contrast, improves substantially — nnU-Net attains $EF\ MAE = 5.52\%$ ($r = 0.828$) against the $\sim 9\%$ originally reported, using $6.5\times$ fewer parameters than the best-Dice model. Protocol modernisation therefore buys clinical agreement rather than contour precision. Second, we identify a reporting error that makes cross-paper boundary comparison unreliable. Computing Hausdorff distances on the network’s resized grid while applying the native image spacing under-reports every distance by the resampling factor — on CAMUS, $\approx 2.2\times$. Under that convention our own models score below 2 mm, which would appear to halve the published baseline; scored at native resolution the same predictions merely match it. Any sub-2 mm HD_{95} reported on CAMUS warrants checking against this failure mode. Code, checkpoints, and per-patient predictions are publicly released.

Keywords: medical image segmentation, echocardiography, CAMUS, deep learning benchmark, U-Net, Swin Transformer, TransUNet

1. Introduction

Two-dimensional echocardiography is the most widely used cardiac imaging modality worldwide: it is portable, inexpensive, non-ionising, and provides real-time visualisation of cardiac structure and function. Its core quantitative output — left ventricular ejection fraction (LVEF), computed from end-diastolic and end-systolic volumes via the biplane Simpson’s method [18] — drives a large fraction of clinical cardiac decisions, from diagnosis of heart failure to monitoring of cardiotoxic chemotherapy. The clinical importance of LVEF is matched by the operator-dependent nature of its measurement: published inter-observer LVEF variability is approximately 7–10% [19], which directly limits the sensitivity of LVEF as a clinical decision boundary.

Automated segmentation of the left ventricular endocardium (LV-endo), epicardium (LV-epi), and left atrium (LA) from two-dimensional apical 2-chamber (2CH) and 4-chamber (4CH) views would, in principle, eliminate this operator-dependent variability and provide reproducible volumetric measurements. The technical challenge is substantial: ultrasound suffers from low signal-to-noise ratio, speckle noise, low boundary contrast, acoustic shadowing, and view-dependent geometric foreshortening.

Architectural pluralism.. Three paradigms have emerged as competitive solutions in this space:

- **Pure convolutional networks.** U-Net [35] and its descendants (residual encoders, SE blocks, attention gates, deep supervision) [28, 29] remain the dominant architectural family, with nnU-Net [11] serving as a strong auto-configured baseline.
- **Hybrid CNN–Transformer networks.** TransUNet [3] and similar hybrids combine a convolutional encoder with a Vision Transformer bottleneck, explicitly trading off local detail (CNN) against global context (self-attention).

*Corresponding author.

Email addresses: faressofiane.mazouni@univ-tebessa.dz (Fares Sofiane Mazouni),
mohamed-yassine.haouam@univ-tebessa.dz (Mohamed Yacine Haouam),
mohamed.amroune@univ-tebessa.dz (Mohamed Amroune), issam.bendib@univ-tebessa.dz (Issam Bendib)

- **Pure Transformer networks.** Swin-UNet [2] replaces the convolutional *backbone* with a hierarchy of windowed self-attention blocks — retaining only a shallow convolutional patch-embedding stem — demonstrating that predominantly attention-based encoders can match or exceed CNN performance on a range of medical tasks.

The benchmarking gap. Despite this architectural pluralism, the literature offers no systematic, like-for-like comparison of these three paradigms on a single shared evaluation pipeline. Each paper introducing a new architecture typically compares against a small number of baselines under different training protocols, augmentations, learning-rate schedules, and evaluation metrics. As a result it is difficult for clinicians or practitioners to answer a basic question: *which architectural family should I use for CAMUS-style apical 2D echocardiographic segmentation, and what are the trade-offs?*

This paper provides that comparison on the CAMUS dataset [19], the standard public 2D echocardiographic segmentation benchmark, with 500 patients across two views and two cardiac phases. We train nine representative architectures spanning the three paradigms above using one identical protocol — same optimiser, loss, schedule, augmentation, mixed-precision settings, and early-stopping criterion — and evaluate them on a single shared pipeline reporting geometric, boundary, clinical, and efficiency metrics simultaneously.

Contributions.

- We present the first unified, like-for-like benchmark of deep segmentation architectures on CAMUS across the three major paradigms: nine architectures spanning pure CNN, hybrid CNN–Transformer, and pure Transformer, trained and evaluated under a single shared protocol.
- We report *multi-dimensional* evaluation, going beyond Dice to include HD95 in millimetres (with per-sample NIfTI pixel spacing), ASSD, biplane Simpson’s EF MAE and Pearson r , ED/ES stratified metrics, robustness across CAMUS image-quality grades, and efficiency benchmarks (parameters, FLOPs, latency, peak memory). The full results matrix and per-patient outputs are released.
- We apply pairwise Wilcoxon signed-rank tests with Bonferroni correction to identify which architectural differences are statistically significant, and produce a critical-difference diagram showing the top-tier band of equivalent models.

- **We document a boundary-metric reporting error that makes published HD95 values on CAMUS mutually incomparable.** Computing the distance transform on the network’s resized 256×256 grid while attaching the native NIfTI spacing under-reports every distance by the resampling factor, which on CAMUS is $\approx 2.2\times$. The effect is large enough to invert a conclusion: under that convention our models score 1.8–2.9 mm and appear to halve the published baseline, while the identical predictions scored at native resolution give 4.1–5.7 mm and merely match it. We report both, and recommend that boundary distances be computed at acquisition resolution — a concrete instance of the validation concerns raised by Reinke et al. [33] and Maier-Hein et al. [26].
- Our best architecture, TransUNet, reaches Dice = 0.9121, HD95 = 4.08 mm; nnU-Net reaches Dice = 0.9099, HD95 = 4.27 mm with $6.5\times$ fewer parameters and roughly $3\times$ faster inference. Boundary accuracy under our unified protocol reproduces the originally published nnU-Net baseline (≈ 4.3 mm) [19] rather than improving on it, whereas EF error falls from $\sim 9\%$ to 5.52%: protocol modernisation shows up in clinical agreement, not in contour precision.
- We release a fully reproducible benchmark framework: training, evaluation, statistical analysis, and figure generation code.

Outline.. Section 2 reviews segmentation architectures and the CAMUS benchmark. Section 3 describes the nine architectures, the unified training protocol, and the evaluation pipeline. Section 4 reports the experimental results. Section 5 discusses architectural design choices, trade-offs, and limitations. Section 6 concludes.

2. Related Work

2.1. Convolutional segmentation networks

U-Net [35] established the encoder–decoder architecture with skip connections that remains the dominant template for medical image segmentation. Subsequent work has refined the encoder (residual blocks / ResNet [8] and later compound-scaled backbones), the decoder (transposed convolution vs. bilinear upsampling, dense connections [10]), and the skip mechanism (attention gates [29], dense skips). Squeeze-and-excitation blocks [9] provide channel-wise recalibration. Atrous spatial pyramid pooling (ASPP) in DeepLabV3+ [4] replaces the deep encoder with dilated convolutions to enlarge the receptive field without down-sampling. Feature Pyramid Networks [21] aggregate features across multiple scales for dense prediction.

The nnU-Net framework [11] integrates these architectural choices with automatic data-dependent hyperparameter selection (target spacing, patch size, batch size, augmentation pipeline) and has set strong baselines on numerous medical benchmarks including CAMUS.

2.2. Transformer-based segmentation networks

Self-attention’s quadratic complexity made global attention impractical for high-resolution images until Vision Transformers [6] demonstrated patch-based tokenisation. TransUNet [3] was the first work to bring this design to medical segmentation, combining a CNN encoder (typically ResNetV2-BiT [17]) with a ViT-B/16 bottleneck and a cascaded upsampler decoder. The hybrid design retains the inductive bias of convolutions at high resolution while granting global context at the bottleneck.

Swin Transformer [23] introduced windowed self-attention with shifted windows, restoring linear complexity in image size. Swin-UNet [2] adapts this to segmentation with a fully Transformer-based U-Net. SwinV2 [22] and Swin-MAE [42] further refine pretraining strategies.

2.3. Echocardiographic segmentation and CAMUS

The CAMUS dataset [19] (*Cardiac Acquisitions for Multi-structure Ultrasound Segmentation*) provides 2D apical-2-chamber and apical-4-chamber acquisitions from 500 patients, each annotated by an expert echocardiographer at end-diastole and end-systole with four classes (background, LV-endo, LV-epi, LA). The official split assigns 400 patients to training, 50 to validation, and 50 to test, with test annotations held out by the challenge organisers. Image quality is graded into *Good*, *Medium*, and *Poor* categories by the same experts.

The original CAMUS paper reported nnU-Net achieving $HD \approx 4.3$ mm on LV-endo and $\sim 9\%$ EF MAE. Subsequent work [31, 36] has reported incremental improvements using deeper architectures, custom losses, and view-aware training. Several papers have additionally explored sequence-aware models on CAMUS half-cycle annotations [31], but these are out of scope for the present benchmark.

Activity on echocardiographic segmentation has continued along four largely separate lines. The first pursues *architectural* refinement: bi-directional token matching between frames [24], dynamic-learning U-Nets that enforce temporal consistency across the cardiac cycle [32], contour-guided query-based fusion aimed explicitly at boundary quality [39], and selective state space

models applied to paediatric echocardiography [43]. The second addresses *domain shift*, which is the dominant obstacle to deploying any of these models across scanners and sites: reinforcement-learning-based domain adaptation [14, 13] and interpretable quantification of the shift itself [7]. The third applies *foundation models* and self-supervision, either by adapting general-purpose promptable segmenters to medical images [25], by generating synthetic cardiac ultrasound for augmentation [38, 34], or by asking what self-supervised echocardiographic representations actually transfer [27]. The fourth targets the *clinical endpoint* directly, predicting or explaining ejection fraction rather than the mask that produces it [44, 37, 20, 41].

What these lines share is a comparison protocol inherited from whichever baseline was convenient. Each reports its own training recipe, its own resolution, and — as we show in Section 4.2 — not always the same definition of a millimetre. That is the gap this benchmark addresses.

EchoNet-Dynamic [30] provides a complementary large-scale video dataset for EF regression. We focus on CAMUS as the standard benchmark for *segmentation-derived* EF. Because EchoNet-Dynamic annotates only the left-ventricular endocardium in a single apical view, it cannot exercise the LV-epicardium, left-atrium, or biplane-EF components of our protocol; whether the architectural trade-offs reported here transfer to other echocardiographic cohorts is therefore a question we leave to future cross-dataset work rather than a claim we make here.

2.4. Validation practice and metric reporting

Running alongside the architectural literature is a smaller body of work arguing that the field’s *measurement* practice, not its model design, is now the limiting factor. Reinke et al. [33] catalogue the ways standard segmentation metrics mislead when applied without regard to structure size, shape, or resolution, and the Metrics Reloaded framework [26] turns that catalogue into problem-driven recommendations for choosing and reporting them. Isensee et al. [12] show that much of the reported progress in 3D medical segmentation dissolves under rigorous validation, with a well-configured U-Net still matching or beating architectures claimed to supersede it. Kazaj et al. [15] reach a similar conclusion for the CNN-versus-Transformer-versus-Mamba question specifically, and Urooj et al. [40] audit a segmentation subfield and find that reported rankings shift once evaluation inconsistencies are removed.

Boundary metrics are especially exposed. Hausdorff distances are reported in millimetres, but a millimetre is only defined once the grid on which the distance transform is computed and the physical spacing attached to that grid refer to the same image. Networks are trained at a

fixed resolution while clinical images are not, so the two can silently diverge. We quantify the consequence on CAMUS in Section 4.2: the same predictions score 1.8 mm or 4.1 mm depending on which convention is used, and only the latter is comparable to the published baseline.

2.5. *What is missing*

A small number of recent works compare a handful of architectures on CAMUS in their experimental sections, but none provide:

- A unified training protocol across CNN, hybrid, and Transformer paradigms (papers typically inherit each architecture’s own published hyperparameters, confounding architecture with protocol).
- Per-class HD95 in millimetres with per-sample NIfTI spacing (most papers report Dice only, or use pixel-space HD95).
- ED- versus ES-stratified reporting (essential for biplane Simpson’s EF analysis).
- Robustness analysis by CAMUS image-quality grade.
- Pairwise Wilcoxon significance with Bonferroni correction and a critical-difference diagram.
- Released model checkpoints and per-patient predictions for independent reanalysis.

This paper aims to fill each of these gaps.

3. Materials and Methods

3.1. *Dataset and splits*

We use the CAMUS dataset [19], comprising 500 patients imaged in apical-2-chamber (2CH) and apical-4-chamber (4CH) views, with expert annotations at end-diastole (ED) and end-systole (ES). Each frame is labelled with four classes: background, LV-endocardium, LV-epicardium, and left atrium. Per-patient pixel spacing is recorded in the NIfTI header and ranges between approximately 0.30 and 0.42 mm.

We adopt the official CAMUS split: 400 patients for training, 50 for validation, and 50 for testing, with no overlap. After ED+ES frame extraction this yields approximately 1,600 training, 200 validation, and 200 test images. Half-cycle sequences are optionally used for augmentation during training; results are reported on the canonical ED+ES test set only.

3.2. *Architectures*

We benchmark nine representative architectures spanning three paradigms:

Pure CNN (seven architectures).

- **UNet-V1** [35]: canonical encoder–decoder with skip connections and trainable transposed-convolution upsampling.
- **UNet-V2** (custom): residual blocks, squeeze-and-excitation [9], attention-gated skips [29], and optional deep supervision.
- **UNet-ResNet**: U-Net with a ResNet-34 encoder [8] pretrained on ImageNet.
- **DeepLabV3+** [4]: ResNet-50 encoder, atrous spatial pyramid pooling bottleneck, lightweight decoder.
- **nnU-Net** [11]: auto-configured U-Net variant, trained from scratch.
- **DenseContextU-Net** (custom): dense blocks [10] with a global-context module in the bottleneck.
- **FPN-UNet**: U-Net decoder with a Feature Pyramid [21] neck over a ResNet-50 encoder.

Hybrid CNN–Transformer (one architecture). **TransUNet** [3]: ResNetV2-BiT encoder [17] feeding a ViT-B/16 bottleneck, with a cascaded upsampler decoder using skip connections from encoder stages.

Pure Transformer (one architecture). **Swin-UNet** [2]: Swin-Tiny encoder [23] pretrained on ImageNet-1k, symmetric Transformer decoder, patch-expanding upsampling.

A complete summary of architectural choices is given in Table 1. We deliberately exclude all SSM-based variants (Mamba, Mamba-2, VMamba) from the present benchmark; those are the subject of a companion study.

3.3. Training protocol

All nine architectures share a single training protocol:

- AdamW optimiser ($\beta_1=0.9$, $\beta_2=0.999$, weight decay 10^{-4}), initial learning rate 10^{-4} , cosine annealing with warm restarts ($T_0=50$ epochs).
- Composite loss: $\mathcal{L} = \mathcal{L}_{\text{CE}} + \mathcal{L}_{\text{Dice}} + 0.1 \mathcal{L}_{\text{boundary}}$, where the boundary term is computed on the signed distance transform [16].
- Image size 256×256 for all CNN and hybrid models; 224×224 for Swin-UNet, which requires input divisible by $\text{window_size} \times 2^{n_{\text{stages}}} = 7 \times 32 = 224$.
- Augmentation: horizontal flip, ± 10 deg rotation, small elastic deformation, mild brightness/contrast jitter.

Table 1: The nine architectures benchmarked in this paper. “Pretrained” indicates whether the encoder is initialised from ImageNet/BiT pretrained weights or trained from scratch.

Paradigm	Architecture	Encoder	Decoder	Pretrained	Params (M)
Pure CNN	UNet-V1	scratch / conv	transposed conv	—	31.04
Pure CNN	UNet-V2	scratch / conv + SE	attn-gated skips	—	33.00
Pure CNN	UNet-ResNet	ResNet-34	U-Net decoder	ImageNet-1k	29.07
Pure CNN	DeepLabV3+	ResNet-50	ASPP + light dec.	ImageNet-1k	40.34
Pure CNN	nnU-Net	auto-configured	deep-supervision	—	15.62
Pure CNN	DenseContextU-Net	dense + context	dense decoder	—	2.18
Pure CNN	FPN-UNet	ResNet-50	FPN + U-Net dec.	ImageNet-1k	32.16
Hybrid	TransUNet	ResNetV2 + ViT-B/16	cascaded upsampler	BiT + IN-21k	102.09
Pure Transformer	Swin-UNet	Swin-Tiny	Swin decoder	IN-1k	41.84

- Up to 100 epochs, automatic mixed precision (AMP, FP16), gradient clipping at $\|\nabla\|_2 \leq 1$ for pretrained-encoder models (Swin-UNet, TransUNet) to prevent the early-epoch divergence we otherwise observed.
- Early stopping with patience 20 on validation mean Dice.
- Training hardware: Google Colab Pro NVIDIA L4 (24 GB) and A100 (40 GB) instances. The final test-set evaluation reported in Section efsec:results was run in a single pass on one NVIDIA RTX 4060 Laptop GPU, so that every metric in every table comes from the same inference run.

All nine architectures were trained with a *uniform* batch size of 32 — the largest value that fits every model within the training-GPU memory budget — rather than a per-architecture batch size, so that the optimisation regime is held as constant as possible across the comparison. We did *not* perform per-architecture hyperparameter search: the goal is explicit cross-architecture comparison under one shared protocol.

3.4. Evaluation metrics

We report a multi-dimensional set of metrics:

Geometric.. Per-class Dice and Intersection-over-Union (IoU) for LV-endo, LV-epi, and LA, computed on the full test set and separately for ED and ES frames.

Boundary.. 95th-percentile Hausdorff distance (HD95) and average symmetric surface distance (ASSD), both in *millimetres* using the per-sample pixel spacing recorded in NIfTI headers. This is consistent with the official CAMUS challenge protocol and is essential for cross-paper comparison.

Clinical.. Biplane Simpson’s ejection fraction (EF) computed from the LV-endo masks at ED and ES across paired 2CH and 4CH views. We report EF MAE (%), Pearson correlation r with the ground-truth EF, and Bland–Altman bias and 95% limits of agreement [1].

Phase-stratified.. All metrics are additionally reported separately for ED and ES frames, which is informative because ED frames tend to be larger and have clearer endocardial boundaries.

Quality-stratified robustness.. Each CAMUS patient carries an expert-assigned image-quality grade (Good / Medium / Poor). A quality-stratified robustness breakdown is deferred to the extended version of this work.

Statistical significance.. For every pair of models we run a Wilcoxon signed-rank test on per-sample mean Dice across the test set, with Bonferroni correction across the $\binom{9}{2} = 36$ pairs. We additionally produce a critical-difference diagram [5] illustrating the band of statistically equivalent top models.

Efficiency.. Trainable parameters, GFLOPs at 256×256 input (measured with `fvcore` on CUDA), single-image inference latency (mean over 100 forward passes after 10 warmup, batch size 1, FP16), and peak GPU memory during forward+backward.

3.5. Reproducibility

A single shared evaluation pipeline is used for all models; the same code path generates the JSON outputs, CSV tables, LaTeX tables, and the figures in this paper. The training pipeline fixes the random seed (42) across PyTorch, NumPy, and Python. Per-model checkpoints, per-test- patient prediction NIfTI files, and per-sample Dice arrays are released to allow independent re-analysis without re-running training.

Table 2: Mean and per-class Dice on the CAMUS official test split for the nine base architectures. The bracketed interval is the bootstrap 95% confidence interval on mean Dice over the 50-patient test set. Best per column in **bold**, second-best underlined. Per-class scores are LV-endocardium, LV-epicardium, and left atrium.

Architecture	Paradigm	Mean Dice	95% CI	LV-endo	LV-epi	LA	Params (M)
TransUNet	Hybrid	0.9121	[0.9078, 0.9162]	0.9367	0.8789	0.9208	102.1
nnU-Net	CNN	<u>0.9099</u>	[0.9054, 0.9143]	0.9371	0.8790	0.9136	15.6
UNet-V1	CNN	0.9085	[0.9041, 0.9130]	0.9369	0.8757	0.9129	31.0
UNet-ResNet	CNN	0.9081	[0.9019, 0.9136]	0.9341	0.8742	0.9160	29.1
UNet-V2	CNN	0.8963	[0.8913, 0.9015]	0.9239	0.8592	0.9059	33.0
Swin-UNet	Transformer	0.8882	[0.8829, 0.8936]	0.9172	0.8520	0.8954	41.8
DeepLabV3+	CNN	0.8602	[0.8511, 0.8681]	0.9109	0.8187	0.8510	40.3
FPN-UNet	CNN	0.8596	[0.8526, 0.8661]	0.8962	0.7944	0.8881	32.2
DenseContextU-Net	CNN	0.8505	[0.8414, 0.8604]	0.8977	0.8046	0.8492	2.2

4. Results

We organise the results into six subsections. Section 4.1 gives the main test-set leaderboard. Sections 4.2 and 4.3 report boundary and clinical (EF) metrics. Section 4.4 examines robustness by image quality. Section 4.5 reports the accuracy $\times$ efficiency trade-off. Section 4.6 reports statistical significance.

4.1. Main test-set Dice

Table 2 summarises mean Dice, per-class Dice, and the parameter count for each architecture on the official CAMUS test split.

Headline ranking. **TransUNet** achieves the highest mean Dice (0.9121), narrowly ahead of **nnU-Net** (0.9099, $\Delta = 0.0022$), **UNet-V1** (0.9085), and **UNet-ResNet** (0.9081). The four top architectures are tightly clustered within $\Delta\text{Dice} \leq 0.004$ — within the Bonferroni-corrected statistical-significance band (Section 4.6), so their ordering should not be read as a ranking. Below this band, **UNet-V2** (0.8963) and **Swin-UNet** (0.8882) form a mid-tier, and **DeepLabV3+** (0.8602), **FPN-UNet** (0.8596), and **DenseContextU-Net** (0.8505) trail.

Table 3: Boundary metrics on the CAMUS test split. HD95 is the 95th-percentile Hausdorff distance, ASSD is the average symmetric surface distance, both reported in millimetres using per-sample NIfTI pixel spacing. Best per column in **bold**.

Architecture	HD95 (mm)	ASSD (mm)
TransUNet	4.08	0.82
UNet-ResNet	4.26	0.84
nnU-Net	4.27	0.84
UNet-V1	4.62	0.87
Swin-UNet	5.08	1.00
UNet-V2	5.29	1.03
DeepLabV3+	5.71	1.19
FPN-UNet	6.21	1.24
DenseContextU-Net	8.29	1.55

Per-class.. LV-endo Dice ranges 0.8962 (FPN-UNet) to 0.9371 (nnU-Net); LA Dice ranges 0.8492 (DenseContextU-Net) to 0.9208 (TransUNet); LV-epi is consistently the hardest class for every architecture (mean across models: Dice $\approx$ 0.85, versus $\approx$ 0.92 for LV-endo).

4.2. Boundary accuracy (HD95 and ASSD in mm)

Boundary metrics are reported in millimetres using the per-sample NIfTI pixel spacing.

HD95.. TransUNet again leads with HD95 = 4.08 mm, followed by UNet-ResNet (4.26 mm) and nnU-Net (4.27 mm), UNet-V1 (4.62 mm), Swin-UNet (5.08 mm), and UNet-V2 (5.29 mm). All boundary distances are computed at the native image resolution: CAMUS frames are approximately 630×519 with a typical spacing of 0.308 mm, so each prediction is resampled back to its native grid before the distance transform is taken. Measuring on the 256×256 network grid while applying the native spacing would under-report every distance by the resampling factor of ~ 2.2 , and we caution that this is an easy error to make silently. The top four architectures sit within 0.2 mm of one another, a spread well inside the training-variance floor discussed in Section 5, and cluster around the published CAMUS nnU-Net baseline of ≈ 4.3 mm [19]. Under a unified modern protocol, then, boundary accuracy on CAMUS appears to have saturated: our reproduction matches the 2019 baseline rather than improving on it. That baseline was in any case

Table 4: End-diastole (ED) and end-systole (ES) stratified mean Dice and HD95 (mm) on the CAMUS test split.

Architecture	Dice		HD95 (mm)	
	ED	ES	ED	ES
nnU-Net	0.9105	0.9093	4.19	4.35
TransUNet	0.9105	0.9137	4.07	4.09
UNet-V1	0.9072	0.9098	4.67	4.57
UNet-ResNet	0.9062	0.9100	4.16	4.37
UNet-V2	0.8989	0.8938	5.06	5.52
Swin-UNet	0.8884	0.8880	5.02	5.14
DeepLabV3+	0.8573	0.8631	5.67	5.76
FPN-UNet	0.8559	0.8633	6.23	6.20
DenseContextU-Net	0.8513	0.8497	7.61	8.98

reported under the original CAMUS evaluation convention, which differs from our per-sample HD95, so the cross-paper comparison is indicative rather than strictly like-for-like.

ED versus ES. Table 4 reports ED- and ES-stratified mean Dice and HD95. The two phases are closely matched for every architecture (mean Dice within ≈ 0.007 and HD95 within ≈ 1.4 mm), with no consistent advantage to either phase. This is a reassuring property for biplane EF, which depends on both phases symmetrically.

4.3. Clinical metric: biplane Simpson’s EF

Table 5 reports biplane Simpson’s EF error per architecture.

Best EF agreement. **nnU-Net** produces the most accurate EF estimates (MAE = 5.52%, $r = 0.828$, Bland–Altman bias +0.84%, 95% limits of agreement $[-14.2, +15.9]\%$, the narrowest of any architecture tested). The near-zero mean bias is comparable to reported echocardiographer inter-observer LVEF error ($\sim 7\text{--}10\%$), but the wide limits of agreement show that per-patient agreement remains looser than the mean error alone suggests; we therefore read this as a mask-derived, like-for-like comparison rather than a claim of clinical parity. UNet-V1 and UNet-ResNet (both MAE = 6.78%) follow, with UNet-ResNet attaining the highest EF correlation ($r = 0.832$, marginally ahead of nnU-Net’s 0.828). Notably, the best-Dice architecture is not

Table 5: Biplane Simpson’s ejection fraction (EF) metrics on the CAMUS test set. Bias and the 95% limits of agreement (LoA) are the Bland–Altman statistics against reference EF over the 50 test patients. Best per column in **bold**.

Architecture	MAE (%)	r	Bias (%)	95% LoA (%)
nnU-Net	5.52	0.828	+0.84	[-14.2, +15.9]
UNet-V1	6.78	0.793	-0.74	[-18.1, +16.6]
UNet-ResNet	6.78	0.832	-2.42	[-18.7, +13.8]
UNet-V2	7.60	0.728	-1.20	[-20.7, +18.3]
TransUNet	7.64	0.799	-3.44	[-21.4, +14.5]
DeepLabV3+	8.46	0.744	-6.10	[-24.0, +11.8]
FPN-UNet	10.12	0.677	-8.12	[-28.6, +12.4]
Swin-UNet	10.46	0.574	-4.10	[-31.0, +22.8]
DenseContextU-Net	16.04	0.444	-11.52	[-48.1, +25.0]

the best on the clinical axis: TransUNet reaches only MAE = 7.64% ($r = 0.799$), behind three CNNs that score lower on bulk overlap. The Transformer-only Swin-UNet shows degraded EF performance (MAE = 10.46%, $r = 0.574$) despite competitive bulk Dice, suggesting it produces accurate *but* biased contours (consistent with its Bland–Altman bias of -4.10%).

4.4. Robustness across image quality

CAMUS assigns each patient an expert image-quality grade (*Good*, *Medium*, *Poor*). Table 6 reports mean Dice stratified by that grade. Every architecture degrades monotonically from Good to Poor, confirming that acquisition quality is a primary driver of segmentation difficulty on CAMUS. The degradation is, however, markedly gentler for the top-tier architectures: TransUNet and nnU-Net lose only ~ 0.02 – 0.03 Dice from Good to Poor and remain above 0.89 on Poor-quality frames, whereas the weakest models (DeepLabV3+, DenseContextU-Net) drop by 0.04 or more and fall below 0.83. Robustness to image quality therefore tracks overall accuracy rather than trading off against it.

4.5. Efficiency and the accuracy–cost trade-off

Table 7 reports parameters, GFLOPs at 256×256 , single-image inference latency, and peak GPU memory.

Table 6: Mean Dice on the CAMUS test split stratified by the expert-assigned image-quality grade (Good $n=94$, Medium $n=76$, Poor $n=30$ test frames). Every architecture degrades monotonically from Good to Poor; the top-tier models remain above 0.88 even on Poor-quality images. Best per column in **bold**.

Architecture	Good	Medium	Poor
TransUNet	0.9186	0.9097	0.8988
nnU-Net	0.9168	0.9091	0.8903
UNet-V1	0.9173	0.9070	0.8848
UNet-ResNet	0.9156	0.9142	0.8734
UNet-V2	0.9091	0.8921	0.8672
Swin-UNet	0.8961	0.8866	0.8675
FPN-UNet	0.8690	0.8628	0.8375
DeepLabV3+	0.8669	0.8555	0.8295
DenseContextU-Net	0.8641	0.8454	0.8209

The accuracy–cost Pareto frontier is dominated at the small-model end by nnU-Net (15.6 M parameters, Dice 0.9099, 2.4 ms latency) and at the high-accuracy end by TransUNet (102.1 M parameters, Dice 0.9121, 7.9 ms). The marginal gain of TransUNet over nnU-Net (+0.0022 Dice — within training variance, Section 4.6) costs $6.5\times$ the parameters and $\approx 3.3\times$ the inference time, making nnU-Net the clear choice for deployment-constrained settings. Figure 1 renders the trade-off as a scatter plot.

Table 7 also exposes a cautionary outlier: at 2.2 M parameters DenseContextU-Net is by far the smallest model, yet it records the highest compute (143.7 GFLOPs) and the slowest inference (25.8 ms) of any architecture in the benchmark. This is a direct consequence of dense connectivity, which repeatedly concatenates and re-convolves feature maps at high spatial resolution; parameter count is therefore a poor proxy for either compute or latency, and the smallest network here is in practice the most expensive to run.

4.6. Statistical significance

We test all $\binom{9}{2} = 36$ pairwise model differences on per-sample mean Dice with the Wilcoxon signed-rank test, Bonferroni-corrected.

Table 7: Efficiency profile on a single NVIDIA L4 GPU. Latency is the mean over 100 forward passes after 10 warmup at batch 1, FP16, 256×256 (or 224×224 for Swin-UNet).

Architecture	Params (M)	FLOPs (G)	Latency (ms)	Peak mem (MB)
DenseContextU-Net	2.2	143.7	25.76	1299
nnU-Net	15.6	23.1	2.38	143
UNet-ResNet	29.1	7.8	2.19	162
UNet-V1	31.0	48.2	1.97	255
FPN-UNet	32.2	205.9	5.48	994
UNet-V2	33.0	52.0	2.73	309
DeepLabV3+	40.3	17.2	2.23	204
Swin-UNet	41.8	13.5	7.09	198
TransUNet	102.1	36.2	7.87	443

Top-tier band.. After Bonferroni correction, the four highest-Dice architectures (TransUNet, nnU-Net, UNet-ResNet, UNet-V1) are not significantly different from one another on bulk mean Dice, while the three lowest (FPN-UNet, DeepLabV3+, DenseContextU-Net) are significantly worse than the top group. Full pairwise p -values are released as supplementary data.

Small Dice gaps are within training variance.. The differences separating the top-tier architectures are of order 0.002–0.005 Dice. We caution against over-interpreting them: in our experiments, re-training an identical architecture under the same nominal protocol in a different compute session shifted its mean Dice by a comparable or larger amount (in the most extreme case, DeepLabV3+ varied by 0.057 between two runs). The bulk-Dice ranking within the top tier should therefore be read as a band of statistically and practically equivalent models, not a strict order. The architecturally meaningful differences in this benchmark are the larger-magnitude ones: the boundary accuracy (HD95), the clinical EF agreement, and the efficiency trade-off, discussed in Section 5.

The critical-difference diagram of Figure 2 renders the banding visually: architectures joined by a horizontal bar are not significantly different in average per-sample rank. Average rank need not track mean Dice. TransUNet attains the best mean Dice yet only a middling average rank (4.09, sixth of nine), because it wins many samples by small margins while occasionally losing by larger ones. It nonetheless falls inside the same top clique as the leading CNNs, which




Figure 1: Accuracy–cost Pareto frontier. Each point is one of the nine architectures, plotted at its (parameter count, mean test Dice). The Pareto front is highlighted.

reinforces that the top tier is a statistically equivalent band rather than a strict order.

4.7. Qualitative results

Figure 3 shows sample predictions from the top four architectures plus a bottom-tier representative on three patients spanning the Good / Medium / Poor image-quality grades.

5. Discussion

5.1. Which architectural family wins on CAMUS?

The top tier of our benchmark is mixed across paradigms: a hybrid CNN–Transformer (TransUNet) leads on bulk Dice and HD95, but a pure CNN (nnU-Net) leads on EF MAE *and* on parameter efficiency. The difference between them on mean Dice is 0.0023, below the Bonferroni-corrected statistical-significance threshold. A practitioner choosing between these two architectures should be guided not by Dice but by deployment constraints (nnU-Net needs 6.5× fewer

Table 8: Pairwise Wilcoxon signed-rank test on per-sample mean Dice, Bonferroni-corrected over $\binom{9}{2} = 36$ pairs. “NS” marks pairs not significantly different at $p < 0.05$ after correction.

	UNet-V1	UNet-V2	UNet-ResNet	DeepLabV3+	nnU-Net	DenseContextU-Net	FPN-UNet	Swin-UNet	TransUNet
UNet-V1	—	2e-13	NS	6e-28	NS	4e-31	2e-30	5e-20	NS
UNet-V2		—	7e-09	4e-16	2e-13	4e-26	7e-21	3e-03	8e-13
UNet-ResNet			—	2e-26	NS	1e-28	1e-29	9e-17	NS
DeepLabV3+				—	3e-28	NS	NS	2e-10	3e-28
nnU-Net					—	6e-32	5e-31	9e-22	NS
DenseContextU-Net						—	NS	4e-22	4e-30
FPN-UNet							—	3e-17	5e-30
Swin-UNet								—	6e-21
TransUNet									—




Figure 2: Critical-difference diagram (average ranks on per-sample Dice, Nemenyi test, $\alpha = 0.05$). Lower average rank is better; a horizontal bar joins architectures whose average ranks differ by less than the critical difference (CD), i.e. that are not significantly different.

parameters and is faster) or by the clinical objective (nnU-Net is more accurate at EF and tighter at Bland–Altman bias).

The pure-Transformer family (Swin-UNet) is competitive on bulk Dice (0.8882) but visibly worse on EF (MAE 10.46%, $r = 0.574$). This finding is consistent with reports that pure-attention models, despite strong overall overlap, can produce boundary contours that are subtly biased [3], an effect that compounds when contours feed Simpson’s biplane volumetric integration.

The bottom of the table is occupied by FPN-UNet, DeepLabV3+, and DenseContextU-Net. All three were designed for natural-image semantic segmentation (FPN/ASPP) or for image classification (dense connectivity) respectively. Their underperformance on CAMUS is consistent with the hypothesis that medical-segmentation-specific designs (U-Net’s skip hierarchy, nnU-Net’s training recipe) matter more on this dataset than generic semantic-segmentation prowess.




Figure 3: Qualitative predictions on three test patients spanning the Good, Medium, and Poor image-quality grades (columns). The top row is the expert ground truth; the remaining rows are, top to bottom, TransUNet, nnU-Net, UNet-ResNet, UNet-V1, and DeepLabV3+. Overlay colours: LV-endocardium (red), LV-epicardium / myocardium (green), and left atrium (blue). Boundary errors concentrate at the LV apex and the mitral-valve plane, and grow visibly from Good to Poor acquisition quality.

5.2. Comparison to the published CAMUS baselines

The original CAMUS paper [19] reported nnU-Net HD95 ≈ 4.3 mm and EF MAE $\sim 9\%$. Our reproduced nnU-Net under the unified protocol reaches HD95 4.27 mm and EF MAE 5.52%. The two axes therefore behave very differently: boundary accuracy is statistically indistinguishable from the 2019 result, while EF error falls by roughly 40%.

We read this as evidence that the contour itself is close to the ceiling imposed by annotation quality on CAMUS. Inter-observer variability on this dataset is of the same order as the differences we measure between architectures — whereas the downstream clinical quantity still has room to improve. Where EF does improve, it is *not* an architectural effect: nnU-Net’s architecture is unchanged. It is attributable to three training-protocol upgrades made standard since 2019:

1. Cosine-with-warm-restarts learning-rate schedule with 10-epoch warmup, in place of the original ReduceLROnPlateau.
2. Mixed-precision training (FP16 AMP), which allows larger effective batch sizes and smoother loss landscapes.
3. Composite loss with an explicit boundary term [16], which directly penalises the Hausdorff-relevant boundary error.

This is a useful baseline calibration point for the field: cross-paper comparisons of CAMUS results from before ~ 2021 to after should be read with the understanding that protocol changes alone produce sub-millimetre HD95 improvements at constant architecture.

5.3. A boundary-metric pitfall, and what it implies for the literature

The correction described in Section 4.2 is worth stating as a general point, because nothing about it is specific to our models. Segmentation networks consume a fixed input size; clinical images do not have one. Any pipeline that resizes an image for inference, scores the prediction on that resized grid, and then multiplies pixel distances by the spacing recorded in the original DICOM or NIfTI header will under-report every boundary distance by the resampling factor. Overlap metrics are immune — Dice and IoU are ratios and cancel the scale — which is precisely why the error survives casual inspection: a pipeline can produce entirely correct Dice and badly wrong HD95 from the same predictions, with no internal inconsistency to notice.

On CAMUS the factor is ≈ 2.2 (native acquisitions average 630×519 against a 256×256 network grid), and the consequence is not cosmetic. It moves our best model from 1.82 mm

to 4.08 mm, and thereby converts “roughly twice as accurate as the published baseline” into “indistinguishable from the published baseline.” Both numbers come from the same weights and the same test split; only the measurement convention differs.

We therefore suggest reading sub-2 mm HD95 values on CAMUS with caution, and we make no claim about which published results are affected — establishing that would require the evaluation code for each, which is usually unavailable. The constructive recommendation is narrower and checkable: report the resolution at which boundary metrics were computed, and prefer acquisition resolution, since that is the only convention under which a millimetre means the same thing across papers. This is a concrete instance of the reporting problems documented by Reinke et al. [33] and Maier-Hein et al. [26], and it is consistent with the broader finding of Isensee et al. [12] and Kazaj et al. [15] that apparent architectural progress in medical segmentation frequently does not survive standardised validation. Our own benchmark is an instance of that pattern rather than an exception to it: once measurement is held fixed, nine architectures spanning three paradigms separate by less than the noise floor on the boundary axis.

5.4. Design choices that matter

Aggregating across our nine architectures, several design choices correlate with strong performance:

- **Pretrained encoders.** Five of the nine architectures initialise their encoder from ImageNet or BiT weights (TransUNet, UNet-ResNet, Swin-UNet, FPN-UNet, DeepLabV3+; Table 1). Pretraining is nonetheless neither necessary nor sufficient for top-tier accuracy on CAMUS: the two scratch-trained models nnU-Net and UNet-V1 place second and fourth on mean Dice — ahead of the pretrained Swin-UNet, FPN-UNet, and DeepLabV3+ — and nnU-Net also attains the best EF agreement at roughly one-sixth of TransUNet’s parameter count. Encoder pretraining thus appears to help mainly where an architecture’s from-scratch inductive bias is weaker, rather than conferring a uniform advantage on this dataset.
- **Deep skip connections** (every U-Net-shaped architecture except DeepLabV3+). DeepLabV3+’s lightweight decoder, while efficient, appears to limit high-resolution boundary precision on CAMUS.
- **Boundary-aware loss** (used uniformly in our protocol). This is a protocol choice rather than an architecture choice. Since our HD95 matches rather than beats the published baseline, we make no claim that it improves absolute boundary accuracy; its role here is to hold the

boundary objective constant across all nine architectures so that the comparison is like-for-like.

- **Auto-configuration** (nnU-Net’s distinguishing feature). The nnU-Net framework produces a competitive model on this dataset with ~ 15 M parameters — substantially smaller than any pretrained-encoder architecture — by selecting patch size, batch size, and augmentation strategy from data statistics.

5.5. Clinical implications

For deployment in echocardiographic LVEF assessment, the choice between the four top-tier architectures should be driven by three criteria: (i) clinical accuracy target, (ii) hardware constraints, and (iii) inference-time latency budget.

- For *maximum reproducibility and tightest EF agreement*, nnU-Net is the recommended choice (lowest EF MAE, smallest Bland–Altman bias).
- For *maximum boundary accuracy* at deployment-irrelevant cost, TransUNet is the highest-HD95 model.
- For *edge / portable ultrasound deployment* where parameters and latency dominate, nnU-Net at 15.6 M parameters remains the best trade-off; no smaller architecture in our benchmark reaches comparable Dice.

None of the four top-tier models reach below-5% EF MAE in this study, indicating that human-level cardiologist agreement is still a target boundary for the field rather than a solved problem.

5.6. Limitations and future work

We acknowledge five limitations.

1. **Single-dataset.** All evaluation is on CAMUS. A cross-dataset transfer evaluation (*e.g.*, EchoNet-Dynamic [30]) would strengthen the generality claims and is planned for future work.
2. **2D only.** CAMUS provides 2D apical views; 3D echocardiography is increasingly common in modern practice and requires architectures and pretraining recipes outside the scope of this benchmark.
3. **Test set size.** The CAMUS official test set is 50 patients (200 ED+ES frames). All point estimates are reported with bootstrap 95% confidence intervals to reflect this.

4. **No SSM models.** The recent wave of selective state space models (Mamba family) is the subject of a separate, dedicated study (a companion study), and is deliberately excluded from the present benchmark to maintain clean three-paradigm framing.
5. **Single training protocol.** A single shared protocol is the goal of this benchmark, but some architectures may benefit from protocol-specific tuning (*e.g.*, larger batch sizes for pretrained Transformers). Per-architecture hyperparameter searches could shift rankings within the top-tier band.

5.7. Take-home messages

1. No single architectural family dominates CAMUS. Hybrid (TransUNet) leads on bulk Dice; pure CNN (nnU-Net) leads on EF accuracy and parameter efficiency. Pure Transformer (Swin-UNet) is competitive on Dice but degraded on EF.
2. Training-protocol modernisation (cosine schedule, AMP, boundary loss) accounts for most of the $\sim 2\times$ improvement in HD95 over the original CAMUS baselines, independent of architecture.
3. LV-epicardium remains the hardest of the three foreground classes across every architectural family.
4. The accuracy–parameter Pareto frontier is dominated by nnU-Net at the efficient end and TransUNet at the high-accuracy end.

6. Conclusion

We have presented the first unified benchmark of deep segmentation architectures on the CAMUS 2D echocardiographic dataset under a single shared training and evaluation protocol. Across nine representative architectures spanning pure CNN, hybrid CNN–Transformer, and pure Transformer paradigms, we report not only mean Dice but also per-class HD95 in millimetres, biplane Simpson’s EF MAE, Bland–Altman bias, ED- and ES-stratified accuracy, image-quality-stratified robustness, parameter / FLOPs / latency efficiency, and Bonferroni-corrected pairwise statistical significance.

The hybrid TransUNet leads on bulk Dice (0.9121) and HD95 (4.08 mm). The pure-CNN nnU-Net leads on EF accuracy (MAE 5.52%, $r = 0.828$) and is $6.5\times$ more parameter-efficient. The top four architectures (TransUNet, nnU-Net, UNet-ResNet, UNet-V1) form a statistically

equivalent top-tier band. The pure-Transformer Swin-UNet is mid-tier on bulk Dice but degraded on EF, suggesting that its contours are subtly biased in a manner that compounds through Simpson’s biplane integration.

Boundary accuracy under the unified protocol (4.08–4.62 mm HD95 across the top four architectures) reproduces the originally reported CAMUS nnU-Net baseline of 4.3 mm rather than improving on it, while EF error falls from the ~9% originally reported to 5.52%. We read this as a useful calibration point for cross-paper comparison: on CAMUS, a decade of protocol modernisation is visible in the downstream clinical quantity but not in the contour itself, which appears to be approaching the ceiling set by annotation variability.

Three directions remain open. (i) Cross-dataset transfer to EchoNet-Dynamic and ACDC would test whether the architectural ranking holds beyond CAMUS. (ii) 3D echocardiography requires architectures and pretraining recipes outside the scope of this benchmark. (iii) The recent wave of selective state space models (Mamba, Mamba-2, VMamba) is the subject of a separate, dedicated companion study (a companion study).

We release the complete training, evaluation, statistical-analysis, and figure-generation pipeline, together with the nine model checkpoints and per-patient JSON / CSV outputs that produced this paper, to lower the barrier to reproducing and extending these results.

Data and code availability

The CAMUS dataset is publicly available from <https://www.creatis.insa-lyon.fr/Challenge/camus/>. Training, evaluation, and analysis code; trained model checkpoints; and all intermediate JSON / CSV outputs supporting this paper are released at <https://github.com/Seraf00/mamba>.

Acknowledgements

The authors thank the CAMUS challenge organisers (CREATIS, INSA Lyon) for making the dataset publicly available. Model training was performed on Google Colab Pro (NVIDIA L4 and A100 GPUs); the final test-set evaluation reported here was run on a single NVIDIA RTX 4060 Laptop GPU. This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Declaration of Generative AI and AI-assisted technologies in the writing process

During the preparation of this work the authors used a large language model to assist with drafting and copy-editing of the manuscript text, and with implementation and refactoring of the released analysis code. After using this tool, the authors reviewed and edited the content as needed and take full responsibility for the content of the published article. No AI-assisted technology was used to generate, alter, or interpret the experimental data or results reported here.

References

- [1] Bland, J.M., Altman, D.G., 1986. Statistical methods for assessing agreement between two methods of clinical measurement. *The Lancet* 327, 307–310.
- [2] Cao, H., Wang, Y., Chen, J., Jiang, D., et al., 2022. Swin-Unet: Unet-like pure transformer for medical image segmentation, in: *ECCV Workshops*.
- [3] Chen, J., Lu, Y., Yu, Q., et al., 2021. TransUNet: Transformers make strong encoders for medical image segmentation. *arXiv preprint arXiv:2102.04306* .
- [4] Chen, L.C., Zhu, Y., Papandreou, G., Schroff, F., Adam, H., 2018. Encoder-decoder with atrous separable convolution for semantic image segmentation, in: *ECCV*, pp. 801–818.
- [5] Demšar, J., 2006. Statistical comparisons of classifiers over multiple data sets. *Journal of Machine Learning Research* 7, 1–30.
- [6] Dosovitskiy, A., Beyer, L., Kolesnikov, A., et al., 2021. An image is worth 16x16 words: Transformers for image recognition at scale, in: *ICLR*.
- [7] Elyasi, S., Serej, N.D., Wall, J., Zolgharni, M., 2026. Domain shift in echocardiography: Interpretable quantification and prediction of cross-dataset left ventricular segmentation. *arXiv preprint arXiv:2607.19643* .
- [8] He, K., Zhang, X., Ren, S., Sun, J., 2016. Deep residual learning for image recognition, in: *CVPR*, pp. 770–778.
- [9] Hu, J., Shen, L., Sun, G., 2018. Squeeze-and-excitation networks, in: *CVPR*, pp. 7132–7141.

- [10] Huang, G., Liu, Z., van der Maaten, L., Weinberger, K.Q., 2017. Densely connected convolutional networks, in: CVPR, pp. 4700–4708.
- [11] Isensee, F., Jaeger, P.F., Kohl, S.A.A., Petersen, J., Maier-Hein, K.H., 2021. nnU-Net: a self-configuring method for deep learning-based biomedical image segmentation. *Nature Methods* 18, 203–211. doi:[10.1038/s41592-020-01008-z](https://doi.org/10.1038/s41592-020-01008-z).
- [12] Isensee, F., Wald, T., Ulrich, C., Baumgartner, M., Roy, S., Maier-Hein, K., Jaeger, P.F., 2024. nnu-net revisited: A call for rigorous validation in 3d medical image segmentation. *arXiv preprint arXiv:2404.09556*.
- [13] Judge, A., Duchateau, N., Judge, T., Sandler, R.A., Sokol, J.Z., Desrosiers, C., Bernard, O., Jodoin, P.M., 2025. Reinforcement learning for unsupervised domain adaptation in spatio-temporal echocardiography segmentation. *arXiv preprint arXiv:2510.14244*.
- [14] Judge, A., Judge, T., Duchateau, N., Sandler, R.A., Sokol, J.Z., Bernard, O., Jodoin, P.M., 2024. Domain adaptation of echocardiography segmentation via reinforcement learning. *arXiv preprint arXiv:2406.17902*.
- [15] Kazaj, P.M., Baj, G., Salimi, Y., Stark, A.W., Valenzuela, W., Siontis, G.C., Zaidi, H., Reyes, M., et al., 2025. From claims to evidence: A unified framework and critical analysis of cnn vs. transformer vs. mamba in medical image segmentation. *arXiv preprint arXiv:2503.01306*.
- [16] Kervadec, H., Bouchtiba, J., Desrosiers, C., Granger, É., Dolz, J., Ben Ayed, I., 2019. Boundary loss for highly unbalanced segmentation, in: MIDL.
- [17] Kolesnikov, A., Beyer, L., Zhai, X., et al., 2020. Big transfer (BiT): General visual representation learning. *ECCV*.
- [18] Lang, R.M., Badano, L.P., Mor-Avi, V., et al., 2015. Recommendations for cardiac chamber quantification by echocardiography in adults. *Journal of the American Society of Echocardiography* 28, 1–39. doi:[10.1016/j.echo.2014.10.003](https://doi.org/10.1016/j.echo.2014.10.003).
- [19] Leclerc, S., Smistad, E., Pedrosa, J., Østvik, A., Cervenansky, F., Espinosa, F., Espeland, T., Berg, E.A.R., Jodoin, P.M., Grenier, T., Lartizien, C., D’hooge, J., Lovstakken, L., Bernard, O., 2019. Deep learning for segmentation using an open large-scale dataset in 2D

- echocardiography. *IEEE Transactions on Medical Imaging* 38, 2198–2210. doi:[10.1109/TMI.2019.2900516](https://doi.org/10.1109/TMI.2019.2900516).
- [20] Li, Y., Zhang, Y., Amadou, A.A., Lai, Y., Zhong, J., Passerini, T., Comaniciu, D., Sharma, P., 2025. Echovlm: Measurement-grounded multimodal learning for echocardiography. arXiv preprint arXiv:2512.12107 .
 - [21] Lin, T.Y., Dollár, P., Girshick, R., He, K., Hariharan, B., Belongie, S., 2017. Feature pyramid networks for object detection, in: *CVPR*, pp. 2117–2125.
 - [22] Liu, Z., Hu, H., Lin, Y., et al., 2022. Swin Transformer V2: Scaling up capacity and resolution, in: *CVPR*.
 - [23] Liu, Z., Lin, Y., Cao, Y., et al., 2021. Swin Transformer: Hierarchical vision transformer using shifted windows, in: *ICCV*.
 - [24] Liu, Z., Tong, L., He, X., Liu, C., Arcucci, R., Jin, C., Zhou, H., 2025. Botm: Echocardiography segmentation via bi-directional optimal token matching. arXiv preprint arXiv:2505.18052 .
 - [25] Ma, J., He, Y., Li, F., Han, L., You, C., Wang, B., 2024. Segment anything in medical images. *Nature Communications* 15 doi:[10.1038/s41467-024-44824-z](https://doi.org/10.1038/s41467-024-44824-z).
 - [26] Maier-Hein, L., Reinke, A., Godau, P., Tizabi, M.D., Buettner, F., Christodoulou, E., Glocker, B., Isensee, F., et al., 2024. Metrics reloaded: Recommendations for image analysis validation. *Nature Methods* doi:[10.1038/s41592-023-02151-z](https://doi.org/10.1038/s41592-023-02151-z).
 - [27] Majchrowska, S., Teare, P., 2026. Evaluating self-supervised echocardiographic representations across downstream extraction strategies for left-ventricular segmentation and ejection fraction estimation. arXiv preprint arXiv:2606.22943 .
 - [28] Milletari, F., Navab, N., Ahmadi, S.A., 2016. V-Net: Fully convolutional neural networks for volumetric medical image segmentation, in: *3DV*.
 - [29] Oktay, O., Schlemper, J., Folgoc, L.L., et al., 2018. Attention U-Net: Learning where to look for the pancreas, in: *MIDL*.

- [30] Ouyang, D., He, B., Ghorbani, A., et al., 2020. Video-based AI for beat-to-beat assessment of cardiac function. *Nature* 580, 252–256. doi:[10.1038/s41586-020-2145-8](https://doi.org/10.1038/s41586-020-2145-8).
- [31] Painchaud, N., Skandarani, Y., Judge, T., et al., 2022. Echocardiography segmentation with enforced temporal consistency. *IEEE Transactions on Medical Imaging* .
- [32] Qu, J., Zhao, J., 2025. A dyl-unet framework based on dynamic learning for temporally consistent echocardiographic segmentation. *arXiv preprint arXiv:2509.19052* .
- [33] Reinke, A., Tizabi, M.D., Sudre, C.H., Eisenmann, M., Rädtsch, T., Baumgartner, M., Acion, L., Antonelli, M., et al., 2021. Common limitations of image processing metrics: A picture story. *arXiv preprint arXiv:2104.05642* .
- [34] Reynaud, H., Gomez, A., Leeson, P., Meng, Q., Kainz, B., 2025. Echoflow: A foundation model for cardiac ultrasound image and video generation. *arXiv preprint arXiv:2503.22357* .
- [35] Ronneberger, O., Fischer, P., Brox, T., 2015. U-Net: Convolutional networks for biomedical image segmentation, in: *MICCAI*, pp. 234–241. doi:[10.1007/978-3-319-24574-4_28](https://doi.org/10.1007/978-3-319-24574-4_28).
- [36] Smistad, E., Østvik, A., Haugen, B., Lovstakken, L., 2021. 2d left ventricle segmentation using deep learning. *IEEE Transactions on Ultrasonics, Ferroelectrics, and Frequency Control* .
- [37] Thomas, S., Cao, Q., Novikova, A., Kulikova, D., Ben-Yosef, G., 2024. Echonnarrator: Generating natural text explanations for ejection fraction predictions. *arXiv preprint arXiv:2410.23744* .
- [38] Tiago, C., Gilbert, A., Beela, A.S., Aase, S.A., Snare, S.R., Sprem, J., 2024. A data augmentation pipeline to generate synthetic labeled datasets of 3d echocardiography images using a gan. *arXiv preprint arXiv:2403.05384* doi:[10.1109/ACCESS.2022.3207177](https://doi.org/10.1109/ACCESS.2022.3207177).
- [39] Ullah, Z., Choi, S., Kim, J., 2026. Contour-guided query-based feature fusion for boundary-aware and generalizable cardiac ultrasound segmentation. *arXiv preprint arXiv:2603.28110* .

- [40] Urooj, A., Abdien, Z.U., Madan, N., 2026. Metrics or mirage? an audit of evaluation inconsistencies in colonoscopy polyp segmentation benchmarks. arXiv preprint arXiv:2607.08203 .
- [41] Wang, Q., He, Z., Liu, Y., Guo, B., Li, Z., Zhao, M., Ju, W., Luo, Z., et al., 2026. Echoagent: Towards reliable echocardiography interpretation with "eyes","hands" and "minds". arXiv preprint arXiv:2604.05541 .
- [42] Xu, Z., Dai, Y., Liu, F., Chen, W., Liu, Y., Shi, L., Liu, S., Zhou, Y., 2023. Swin MAE: Masked autoencoders for small datasets. Computers in Biology and Medicine 161, 107037. doi:[10.1016/j.compbiomed.2023.107037](https://doi.org/10.1016/j.compbiomed.2023.107037).
- [43] Ye, Z., Chen, T., Wang, F., Zhang, H., Zhang, L., 2024. P-mamba: Marrying perona malik diffusion with mamba for efficient pediatric echocardiographic left ventricular segmentation. arXiv preprint arXiv:2402.08506 .
- [44] Zhang, Y., Liu, B., Bunting, K.V., Brind, D., Thorley, A., Karwath, A., Lu, W., Zhou, D., et al., 2024. Development of automated neural network prediction for echocardiographic left ventricular ejection fraction. arXiv preprint arXiv:2403.12152 .